

[blog.codinghorror.com](http://blog.codinghorror.com)

# A Visual Explanation of SQL Joins

*Jeff Atwood*

4 minutes

11 Oct 2007

I thought Ligaya Turmelle's [post on SQL joins](#) was a great primer for novice developers. Since SQL joins *appear* to be set-based, the use of [Venn diagrams](#) to explain them seems, at first blush, to be a natural fit. However, like the commenters to her post, I found that the Venn diagrams didn't quite match the [SQL join syntax](#) reality in my testing.


I love the concept, though, so let's see if we can make it work. Assume we have the following two tables. **Table A** is on the left, and **Table B** is on the right. We'll populate them with four records each.

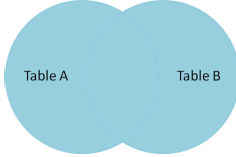
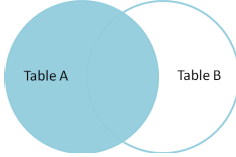
id	name	id	name
1	Pirate	1	Rutabaga
2	Monkey	2	Pirate
3	Ninja	3	Darth Vader
4	Spaghetti	4	Ninja

Let's join these tables by the name field in a few different ways and see if we can get a conceptual match to those nifty Venn diagrams.

```
SELECT * FROM TableA
INNER JOIN TableB
ON TableA.name = TableB.name
id name    id  name
```

1	Pirate	2	Pirate
3	Ninja	4	Ninja



<p><b>Inner join</b> produces only the set of records that match in both Table A and Table B.</p>																													
<p>SELECT * FROM TableA</p> <p><b>FULL OUTER JOIN</b> TableB</p> <p>ON TableA.name = TableB.name</p> <table><tr><td>id</td><td>name</td><td>id</td><td>name</td></tr><tr><td>1</td><td>Pirate</td><td>2</td><td>Pirate</td></tr><tr><td>2</td><td>Monkey</td><td>null</td><td>null</td></tr><tr><td>3</td><td>Ninja</td><td>4</td><td>Ninja</td></tr><tr><td>4</td><td>Spaghetti</td><td>null</td><td>null</td></tr><tr><td>null</td><td>null</td><td>1</td><td>Rutabaga</td></tr><tr><td>null</td><td>null</td><td>3</td><td>Darth Vader</td></tr></table> <p><b>Full outer join</b> produces the set of all records in Table A and Table B, with matching records from both sides where available. If there is no match, the missing side will contain null.</p>	id	name	id	name	1	Pirate	2	Pirate	2	Monkey	null	null	3	Ninja	4	Ninja	4	Spaghetti	null	null	null	null	1	Rutabaga	null	null	3	Darth Vader	
id	name	id	name																										
1	Pirate	2	Pirate																										
2	Monkey	null	null																										
3	Ninja	4	Ninja																										
4	Spaghetti	null	null																										
null	null	1	Rutabaga																										
null	null	3	Darth Vader																										
<p>SELECT * FROM TableA</p> <p><b>LEFT OUTER JOIN</b> TableB</p> <p>ON TableA.name = TableB.name</p> <table><tr><td>id</td><td>name</td><td>id</td><td>name</td></tr><tr><td>--</td><td>----</td><td>--</td><td>----</td></tr><tr><td>1</td><td>Pirate</td><td>2</td><td>Pirate</td></tr><tr><td>2</td><td>Monkey</td><td>null</td><td>null</td></tr><tr><td>3</td><td>Ninja</td><td>4</td><td>Ninja</td></tr><tr><td>4</td><td>Spaghetti</td><td>null</td><td>null</td></tr></table> <p><b>Left outer join</b> produces a complete set of records from Table A, with the matching records (where available) in Table B. If there is no match, the right side will contain null.</p>	id	name	id	name	--	----	--	----	1	Pirate	2	Pirate	2	Monkey	null	null	3	Ninja	4	Ninja	4	Spaghetti	null	null					
id	name	id	name																										
--	----	--	----																										
1	Pirate	2	Pirate																										
2	Monkey	null	null																										
3	Ninja	4	Ninja																										
4	Spaghetti	null	null																										


```
SELECT * FROM TableA
LEFT OUTER JOIN TableB
ON TableA.name = TableB.name
WHERE TableB.id IS null
id name    id  name
```

---

```
2  Monkey  null null
```

```
4  Spaghetti null null
```

To produce the set of records only in Table A, but not in Table B, we perform the same left outer join, then **exclude the records we don't want from the right side via a where clause.**



```
SELECT * FROM TableA
FULL OUTER JOIN TableB
ON TableA.name = TableB.name
WHERE TableA.id IS null
OR TableB.id IS null
id name    id  name
```

---


```
2  Monkey  null null
```

```
4  Spaghetti null null
```

```
null null    1  Rutabaga
```

```
null null    3  Darth Vader
```

To produce the set of records unique to Table A and Table B, we perform the same full outer join, then **exclude the records we don't want from both sides via a where clause.**



There's also a cartesian product or **cross join**, which as far as I can tell, can't

be expressed as a Venn diagram:

```
SELECT * FROM TableA
```

```
CROSS JOIN TableB
```

This joins "everything to everything", resulting in  $4 \times 4 = 16$  rows, far more than we had in the original sets. If you do the math, you can see why this is a *very* dangerous join to run against large tables.