

## Part 1 : Micro-services



Follow

Dec 20, 2018 ·  Unlisted

Hike



For any product or service, availability is one of the most important factors and when you are a messaging app it becomes even more important to have highest availability of your services with lowest latency.

*“Availability is not just sending a response back, it has to be in a specified time.”*

We had an interesting journey from where we started to even think of four nines of availability for all our services.

Just to give a glimpse of what it means:

- Five-nines or 99.999% availability means 5 minutes, 15 seconds or less of downtime in a year
- Four nines or 99.99% availability allows 52 minutes, 36 seconds downtime per year
- Three nines or 99.9% availability allows 8 hours, 46 minutes downtime per year

- *Two nines or 99% availability allows 3 days, 15 hours and 40 minutes downtime per year.*
- *One nine or 9% availability allows over 332 days of downtime per year.*

At the first glance, 52 minutes of downtime in a year seems too much but what this means is that you can vary your SLA's for just 8.65 seconds in an entire day (including peak traffic hours)

## Why is five, four or even three nines of availability so hard to achieve?


1. *Bugs: Human error still accounts to around 60% of the downtime*
2. *SPOF: Single point of failures leading to system wide downtime*
3. *Hardware failures, Network issues*
4. *Unusual traffic spikes*
5. *Third party or external APIs dependencies*

In order to improve on availability of overall system, you need to work on each of the above points — after all the “*Strength of a chain is as good as strength of its weakest link*”.

## So, where do you start if you want to improve on availability?

### 1. Monitoring


The first and foremost step to improve on anything is to monitor it regularly and have clear alerts whenever it fails




### 2. Eradicate Single point of failures

Back in 2012 when we started, we had a monolith architecture where all the components were together in a single deployable unit and all the processing was done synchronously.

It was literally something like this.






Though it gave us the speed (initially), as it was quick to develop but there were many pitfalls like:

- Single point of failures (SPOF), so if there is an issue in one of the flows, the whole system gets affected.
- Difficult to do fault detection and isolation
- Continuous deployment become an issue
- Application becomes large and too complex to understand

To avoid these pitfalls, we broke down our architecture to a service-oriented one with clear separation of services, which not only helped to eradicate SPOF but also make fault detection and isolation easy.



Now, let's say 'Db Y' is down then only service Y or Processor 1 will get affected and rest of the system will remain healthy. For fault detection and isolation we are using Hystrix for some and an internal tool for other layers.

### 3. Tackling Bugs

*"Monitor → Fix → Automate → Repeat"*

Monitor as many things you can, run your automation test suite as early as possible, make hard checks at every step to stop bad code/configuration from reaching your production environment.

## Stop bad code from reaching the production system

Bugs often reach the production systems because of missing edge condition, developer didn't get it reviewed or just missed to run test cases around it.

To prevent this, we have put the following mandatory checks on every 'pull request' raised by a developer.

- All automated test suites must have run successfully
- Should be reviewed by at least one other developer

### [MI-2609] IOException handling #6562

[Edit](#)

**Open** pankajsinghal wants to merge 1 commit into `master` from `cleaning`

Conversation 0 Commits 1 Checks 0 Files changed 16 +1 -19

pankajsinghal commented 16 days ago

No description provided.

removed unnecessary IOException throw ✗ 331c3db

pankajsinghal requested review from ashishmjn and aditya-bsb 16 days ago

pankajsinghal changed the title from removed unnecessary IOException throw to [MI-2609] IOException handling a minute ago

pankajsinghal assigned aditya-bsb a minute ago

pankajsinghal added the `TestRequest` label just now

Add more commits by pushing to the `cleaning` branch on `hike/hike-mq`.

**Review required** [Show all reviewers](#)

At least 1 approving review is required by reviewers with write access. [Learn more](#).

**All checks have failed** [Hide all checks](#)

2 failing checks

**default** — Build finished. [Details](#)

**tests** — Hike-Jenkins Required [Details](#)

**Merging is blocked**

Merging can be performed automatically with 1 approving review.

[Merge pull request](#) You can also [open this in GitHub Desktop](#) or view [command line instructions](#).

**Reviewers**

aditya-bsb ✓

ashishmjn ●

At least 1 approving review is required to merge this pull request.

**Assignees**

aditya-bsb

**Labels**

`TestRequest`

**Projects**

None yet

**Milestone**

No milestone

**Notifications**

[Unsubscribe](#)

You're receiving notifications because you authored the thread.

**2 participants**

[Lock conversation](#)

### Pull Request (PR) Raised

### [MI-2609] IOException handling #6562

[Edit](#)

**Open** pankajsinghal wants to merge 1 commit into `master` from `cleaning`

Conversation 0 Commits 1 Checks 0 Files changed 16 +1 -19

pankajsinghal commented 16 days ago

No description provided.

removed unnecessary IOException throw ✗ 331c3db

**Reviewers**

aditya-bsb ✓

ashishmjn ●

**Assignees**

aditya-bsb

**PR Approved**

## Edit


<https://blog.hike.in/4-9s-b1fc497f3de2>

**PR Passed**

#### 4. Taking code to production

Last but not the least, is how you take code to production system.

- Do a Canary deployment. Route small traffic to the new code base and slowly increase traffic to 100%.



- Monitoring and Alerting
  - ➔ For monitoring, we have build in-house tools over [Grafana](#) and BigQuery for tracking both tech and product issues. Also, we use [Overops](#) for exception detection.
  - ➔ For alerting we are using Nagios and VictorOps.

In next blog of “Availability series”, we will deep dive into specific components of our messaging architecture and how we improved on availability.